


# E96 MNIST Final Project: Single Digit Image Classifier

Steven Bash and Jacob Apoian



# Design Process

- Started without the skeleton
- We figured out
  - How to import libraries
  - How to create neural network class and forward functions
  - How to import the training/testing data
  - How to create our train/test sets
  - Instantiate model and train it
  - We decided to use Adam as the optimizer

And we achieved 10% accuracy 🎉

# Design Process


- Then we tried merging our with the skeleton
  - Ran into bugs
  - But, we figured out why we got 10% accuracy
- In the training loop, we wrote `model.zero_grad()` instead of `optimizer.zero_grad()`
  - Essentially, although we had chosen an optimizer, we weren't training our model
- So, we continued to work on our original code and didn't end up using the skeleton
- Then, we wanted to create an even better model
  - For all tests, we used 10 epochs and trained with the first 60000 of the MNIST data set

Iterations

## 1. All FC NN with Adam using F.nll\_loss()

$$9744/10000 = .9744$$

We started with 4 fully connected layers and using the Nll\_loss loss function



2. Then we changed the loss function to `nn.CrossEntropyLoss()`

Still All FC NN with Adam

$9768/10000 = .9768$

tensor(0.0110, grad\_fn=<NllLossBackward0>)

tensor(2.2650e-07, grad\_fn=<NllLossBackward0>)

tensor(0.0916, grad\_fn=<NllLossBackward0>)

tensor(0.0396, grad\_fn=<NllLossBackward0>)

tensor(0.0017, grad\_fn=<NllLossBackward0>)


tensor(0.6452, grad\_fn=<NllLossBackward0>)

tensor(0.0002, grad\_fn=<NllLossBackward0>)

tensor(0.0670, grad\_fn=<NllLossBackward0>)

tensor(1.1921e-08, grad\_fn=<NllLossBackward0>)

tensor(4.5299e-07, grad\_fn=<NllLossBackward0>)



Trying we tried to get fancy!

Maxpool originally 2, increased to 14 (lowered dimension from 9216 to 64)

Running at 10 epochs with increased maxpool took:

1334.447 seconds, 22 minutes

Originally, a single epoch took:  $\sim 5 \text{ minutes} * 10 \text{ iterations} = 50 \text{ minutes}$

```
... tensor(0.0122, grad_fn=<NllLossBackward0>)
```




### 3. Forward function with 2 FC, 2 CL, and 2 dropout layers

8118/10000 = 0.8118

48424/60000 = 0.8070666666666667

tensor(1.1661, grad\_fn=<NllLossBackward0>)  
tensor(0.7710, grad\_fn=<NllLossBackward0>)  
tensor(0.2555, grad\_fn=<NllLossBackward0>)  
tensor(0.9164, grad\_fn=<NllLossBackward0>)  
tensor(0.6896, grad\_fn=<NllLossBackward0>)  
tensor(0.3064, grad\_fn=<NllLossBackward0>)  
tensor(0.6938, grad\_fn=<NllLossBackward0>)  
tensor(0.2539, grad\_fn=<NllLossBackward0>)  
tensor(0.8421, grad\_fn=<NllLossBackward0>)  
tensor(0.4648, grad\_fn=<NllLossBackward0>)



# Design Process


- Through trial and error we changed the amount and type of layers used
  - We didn't record data for every trial
- We ran into many issues with matching the shape of the forward function with the shape of the model
  - Unsuccessfully tried to add a single CL and a pool layer
  - Successfully tried a dropout layer with .5 rate

## 4. FC layers and dropout layer set to .5

$$9650/10000 = 0.965$$


$$58712/60000 = 0.9785333333333334$$

tensor(0.0906, grad\_fn=<NllLossBackward0>)  
tensor(0.0393, grad\_fn=<NllLossBackward0>)  
tensor(0.0035, grad\_fn=<NllLossBackward0>)  
tensor(0.0037, grad\_fn=<NllLossBackward0>)  
tensor(0.5758, grad\_fn=<NllLossBackward0>)  
tensor(0.0310, grad\_fn=<NllLossBackward0>)  
tensor(0.0160, grad\_fn=<NllLossBackward0>)  
tensor(0.0547, grad\_fn=<NllLossBackward0>)  
tensor(0.5861, grad\_fn=<NllLossBackward0>)  
tensor(0.0102, grad\_fn=<NllLossBackward0>)



# Final Result

- We chose our most accurate model
  - Testing accuracy: 97.68%
  - Training accuracy. We didn't save it :/
- 
- Using FC NN only & Adam optimizer
  - Using `nn.CrossEntropyLoss()` loss function



# Best vs Worst

## 97.68% accuracy - fully connected only

tensor(0.0110, grad\_fn=<NiILossBackward0>)

tensor(2.2650e-07, grad\_fn=<NiILossBackward0>)

tensor(0.0916, grad\_fn=<NiILossBackward0>)

tensor(0.0396, grad\_fn=<NiILossBackward0>)

tensor(0.0017, grad\_fn=<NiILossBackward0>)

tensor(0.6452, grad\_fn=<NiILossBackward0>)

tensor(0.0002, grad\_fn=<NiILossBackward0>)

tensor(0.0670, grad\_fn=<NiILossBackward0>)

tensor(1.1921e-08, grad\_fn=<NiILossBackward0>)

tensor(4.5299e-07, grad\_fn=<NiILossBackward0>)

## 81.18% accuracy - FC, convolution and dropout layers

tensor(1.1661, grad\_fn=<NiILossBackward0>)

tensor(0.7710, grad\_fn=<NiILossBackward0>)

tensor(0.2555, grad\_fn=<NiILossBackward0>)

tensor(0.9164, grad\_fn=<NiILossBackward0>)

tensor(0.6896, grad\_fn=<NiILossBackward0>)

tensor(0.3064, grad\_fn=<NiILossBackward0>)


tensor(0.6938, grad\_fn=<NiILossBackward0>)

tensor(0.2539, grad\_fn=<NiILossBackward0>)


tensor(0.8421, grad\_fn=<NiILossBackward0>)

tensor(0.4648, grad\_fn=<NiILossBackward0>)

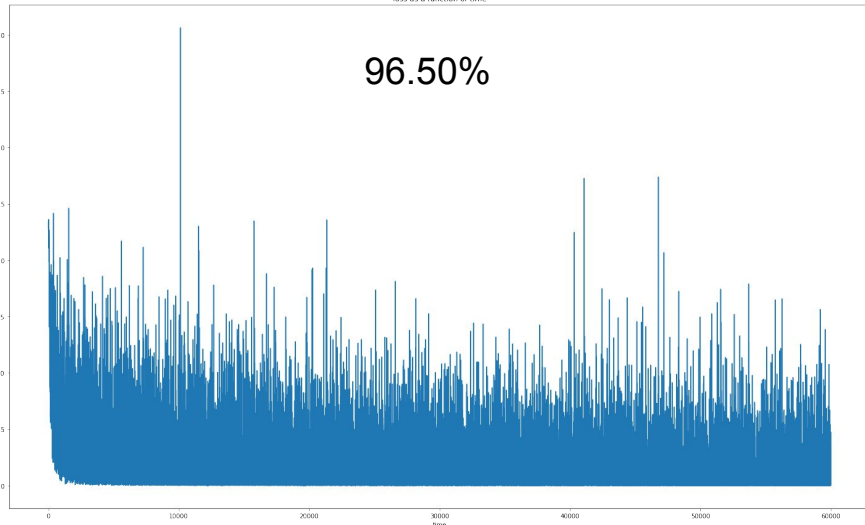
Best 97.68%



97.44%



96.50%



Worst 81.18%

